

# ECE 219 Project 2

Tanner Dulay, Loïc Maxwell, Zoë Tucker

February 3, 2021

## 1 Question 1

In this assignment, we explore various methods for clustering of text data. To begin, we use the “20 Newsgroups” dataset, taking the two well-separated classes previously defined (“computer technology” and “recreational activity”). As this project uses unsupervised learning, we do not distinguish between training and test data, and instead use the entire dataset throughout the project.

To start our process, we build the TF-IDF matrix. We use `min_df = 3`, exclude English stopwords/punctuation, and remove headers and footers, but we do not perform stemming or lemmatization. The resulting TF-IDF matrix is  $7882 \times 21909$ .

## 2 Question 2

Next, we apply k-means clustering directly to the TF-IDF data. As we are interested in two classes of documents, we set  $k = 2$ . The below contingency matrix (Figure 1) shows the results of this.




Figure 1: Contingency Matrix for K-Means on TF-IDF Matrix

### 3 Question 3

We then find the metrics we need to evaluate our classification results. We find the following values:

|                                   |                    |
|-----------------------------------|--------------------|
| Homogeneity score                 | 0.5600025412974119 |
| Completeness score                | 0.5716557539454379 |
| V_measure score                   | 0.5657691483417876 |
| Adjusted Rand Index score         | 0.6327475635535726 |
| Adjusted Mutual Information Score | 0.5657289846874574 |

### 4 Question 4

As the metrics found in Question 3 do not demonstrate a good clustering result, we now try dimensionality reduction to help our data better fit the assumptions made by k-means algorithms.

We begin by finding the percent of variance the top  $r$  principal components can retain for various values of  $r$ . Figure 2 shows our findings.




Figure 2: Principal Component Results

## 5 Question 5

Next, we will try to experimentally find the value of  $r$  (i.e. the number of dimensions) that we can reduce our data to in order to obtain the best clustering results. Testing  $r \in \{1, 2, 3, 5, 10, 20, 50, 100, 300\}$  using both NMF and SVD, we plot the values of the five metrics found in Question 3 for each  $r$ . The results are shown in Figures 3 and 4.




Figure 3: NMF Results




Figure 4: SVD Results

## 6 Question 6

As we can see, the behavior of the measures as  $r$  increases is non-monotonic.

This is due to a trade-off between a larger value of  $r$  giving a more faithful representation of the data, and a smaller value of  $r$  being better suited for k-means clustering. The higher the dimensionality of a dataset, the more difficult it is to meaningfully evaluate the “distance” between data points, which is what k-means relies on. Thus, we expect there to be an optimal value of  $r$  that balances these two concerns and maximizes our scores. In our case, we find that  $r = 2$  for NMF, and  $r = 50$  for SVD.

## 7 Question 7

Now, we want to visualize our clustering results. We do this by using k-means clustering with our dataset reduced by NMF and SVD with the optimal values of  $r$  found in Question 6, then projecting the data into 2-dimensional space with SVD and coloring the points according to their label. Figures 5 and 6 show the results for NMF with  $r = 2$ . Figures 7 and 8 show the results for SVD with  $r = 50$ .




Figure 5: Predicted labels for NMF reduced data




Figure 6: Ground truth labels for NMF reduced data




Figure 7: Predicted labels for SVD reduced data




Figure 8: Ground truth labels for SVD reduced data

## 8 Question 8

Because the NMF reduced data was already in 2D, we did not reduce it further with SVD. However, if we were to do SVD reduction, the resulting matrix would still be positive, as it is in Figures 5 and 6. SVD is an orthogonal projection that preserves the angles between feature vectors.

The 2D visualization from Figure 8 shows that the SVD reduced data looks somewhat linearly separable around a value of 0 for Dimension 1, which stems from the second principal component of the TF-IDF matrix. This suggests that this principal component holds important information pertaining to the class that the data belongs to. Although this analysis is restricted to two dimensions, this also suggests that this type of distribution would be good for KMeans clustering. A similar type of separation can be found at the diagonal between dimensions 0 and 1 for the NMF reduced data in Figure 6. This would also make it ideal for KMeans clustering.

However, for both types of reductions, because there is no way to perfectly linearly separate the data there are some better options out there for properly classifying this data. Perhaps a classifier that can learn non-linear behavior in the data, such as a neural network with ReLU/eLU, can perform better on this dataset.

## 9 Question 9

Using SVD with  $r=50$ , which gave us the best results on the two classes, we got the following contingency matrix for KMeans on the entire dataset of 20 classes:




Figure 9: Contingency Matrix for KMeans on 20 Class Data

The respective scores for this clustering are as follows:

|                                   |                     |
|-----------------------------------|---------------------|
| Homogeneity score                 | 0.31630458145750995 |
| Completeness score                | 0.3682552023183464  |
| V_measure score                   | 0.3403086491478565  |
| Adjusted Rand Index score         | 0.10379799750286293 |
| Adjusted Mutual Information Score | 0.3380068278889245  |



## 10 Question 10

Using the Kullback-Leibler Divergence loss metric for NMF and  $r=10$ , we were able to get the following scores:

|                                   |                     |
|-----------------------------------|---------------------|
| Homogeneity score                 | 0.37714629232162333 |
| Completeness score                | 0.3900505160910544  |
| V_measure score                   | 0.3834898799077068  |
| Adjusted Rand Index score         | 0.18833101982676798 |
| Adjusted Mutual Information Score | 0.38146535004967785 |

There is indeed a noticeable improvement over the performance using ordinary SVD with  $r=50$  (from Q9).

## 11 Question 11

We compared the performance of KMeans on the dataset reduced with UMAP using various values for  $n\_components$ . The plots are shown below:




Figure 10: Metric scores for Cosine Metric vs.  $n\_components$




Figure 11: Metric scores for Euclidean Metric vs. n\_components

The best scores using both metrics came from **n\_components=5**:

| Cosine Metric                     |                    |
|-----------------------------------|--------------------|
| Homogeneity score                 | 0.5347926404764849 |
| Completeness score                | 0.5549487009354954 |
| V_measure score                   | 0.5446842655666289 |
| Adjusted Rand Index score         | 0.4167210932677909 |
| Adjusted Mutual Information Score | 0.5431756896373161 |

| Euclidean metric                  |                       |
|-----------------------------------|-----------------------|
| Homogeneity score                 | 0.008428056566878686  |
| Completeness score                | 0.008490790660743977  |
| V_measure score                   | 0.008459307306639803  |
| Adjusted Rand Index score         | 0.0014762057832842827 |
| Adjusted Mutual Information Score | 0.00524781170211466   |

## 12 Question 12

Using  $n\_components = 5$  as determined above, we get the below contingency matrices – Figure 12 corresponds to the Euclidean metric, and Figure 13 corresponds to the cosine metric.




Figure 12: Confusion matrix for UMAP-reduced Kmeans with Euclidean metric




Figure 13: Confusion matrix for UMAP-reduced Kmeans with cosine metric

The cosine matrix is obviously much better, which corresponds with the greatly improved scores using this metric.

Examining the cosine metric more closely, we see that most categories were well identified, but some were more likely to be confused with each other. Tracing the connections, we can see that points with ground truth label 19 were likely to be identified with points labeled 0 or 15, and points from label 12 were misidentified as 1 or 6. In the original dataset, this means that points from `talk.religion.misc` were identified as `alt.atheism` or `soc.religion.christian`, and points from `sci.electronics` were identified as `comp.graphics` or `misc.forsale`. The religion groups form an intuitive grouping, and the confusion between the latter three groups perhaps indicates that selling hardware and electronics was

common on these newsgroups.

## 13 Question 13

Using the best hyperparameter combination from Q11/12, `n_components=5` and the cosine metric, we reduced the Tfidf matrix with UMAP. The scores for agglomerative clustering using 'ward' and 'single' linkage on this data are shown below:

| Ward linkage                      |                    |
|-----------------------------------|--------------------|
| Homogeneity score                 | 0.524473300498607  |
| Completeness score                | 0.5479912976714623 |
| V_measure score                   | 0.53597443687123   |
| Adjusted Rand Index score         | 0.3940501791969448 |
| Adjusted Mutual Information Score | 0.5344316766896055 |

| Single linkage                    |                        |
|-----------------------------------|------------------------|
| Homogeneity score                 | 0.022917457355135236   |
| Completeness score                | 0.3464156710750703     |
| V_measure score                   | 0.04299081646294147    |
| Adjusted Rand Index score         | 0.00046761208355549543 |
| Adjusted Mutual Information Score | 0.03748616963750802    |

It's clear that ward linkage performs significantly better on the dataset. Single linkage uses the minimum distance between all points in each cluster to determine whether to merge them or not, so it's a lot more sensitive to outliers and tends to result in uneven cluster sizes. Ward linkage on the other hand, minimizes the variance between pairs of clusters, which can produce more regular cluster sizes and hence better results for our particular dataset.

## 14 Question 14

In this section we investigate the DBSCAN and HDBSCAN clustering algorithms. We kept UMAP with `n_components=5` and the cosine metric as the dimensionality reduction algorithm. The HDBSCAN `min_cluster_size` was set to 100, and other hyperparameters were tuned to improve the performance of each of these algorithms. We performed hyperparameter tuning on DBSCAN and HDBSCAN for epsilon values of [0.01,0.05,0.1,0.3,0.5,0.7,0.9,1], metric types of [Euclidean, cosine, Manhattan] (if applicable), and minimum sample values of [10, 20, 50, 70, 100]. The following table depicts the scores for each clustering algorithm for the hyperparameter combination that resulted in the best average score across all 5 clustering metrics:

| <b>DBSCAN: Epsilon = 0.5, Metric = Manhattan, Minimum samples = 10</b> |                     |
|--|---------------------|
| Homogeneity score  | 0.5407607595568981  |
| Completeness score   | 0.45471096739764455 |
| V_measure score  | 0.49401673890037173 |
| Adjusted Rand Index score  | 0.2967260971944622  |
| Adjusted Mutual Information Score                                      | 0.47920988791262803 |

| <b>HDBSCAN: Epsilon = 0.3, Metric = Euclidean, Minimum samples = 10</b> |                     |
|---|---------------------|
| Homogeneity score   | 0.43889528158878366 |
| Completeness score  | 0.5288819318072416  |
| V_measure score   | 0.47970500064413385 |
| Adjusted Rand Index score   | 0.2565568616115164  |
| Adjusted Mutual Information Score                                       | 0.4780526456871693  |

## 15 Question 15

DBSCAN was the better performing clustering algorithm. The associated contingency matrix is shown below:




Figure 14: Contingency Matrix for DBSCAN on 20 Class Data

The contingency matrix indicates that there must be at least 96 clusters since that's the largest number shown on the x-axis ticks. By looking at the number of unique labels in the trained DBSCAN model, we are able to determine that there are actually 168 clusters formed by the algorithm (excluding the -1 labels). This number is high because of our low value for the minimum samples hyperparameter, so most of these clusters contain few data points. However, most of the data points were captured by the larger clusters shown in the contingency matrix, which resulted in acceptable clustering metric scores.

A label of -1 indicates that the data point was classified as noise, and was not assigned to any cluster. This cluster label is represented as cluster number 0 in the contingency matrix. We can tell because this cluster contains many

data points from each of the 20 classes.

## 16 Question 16

We used the “read\_csv” function of the Pandas library to acquire the BBC dataset as a dataframe. After converting the written class labels to numerical target values for each document, we then used our established pipeline to vectorize (excluding stopwords/punctuation) then normalize the data and analyze the dimensions of the resulting TF-IDF matrix. As with the 20 Newsgroups dataset, we then reduced dimensionality with UMAP and searched for the optimal parameter configuration for multiple clustering algorithms. Starting with KMeans, we tested how our five aforementioned accuracy metrics vary with the number of principle components used in clustering for both Euclidean (Figure 15) and Cosine distance metrics (Figure 16). The best combination was  $n\_components = 10$  with the Euclidean distance metric.




Figure 15: Variation in accuracy metrics by number of components for Kmeans Clustering and Euclidean distance



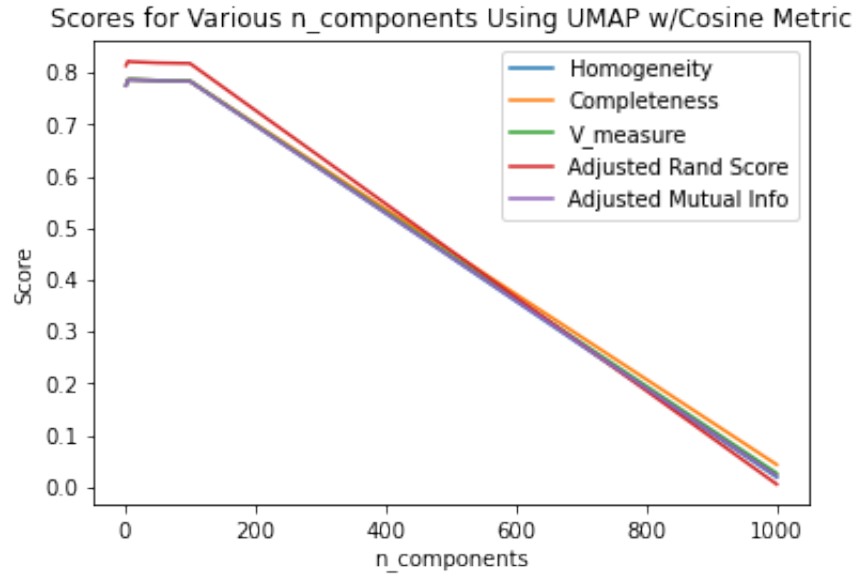


Figure 16: Variation in accuracy metrics by number of components for Kmeans Clustering and Cosine distance

We then tested the accuracy of the DBSCAN clustering algorithm for both distance metrics. We cycled through values from .01 to 1 for epsilon and from 1 to 50 for the minimum number of points per cluster to find the best combination of hyperparameters for our analysis. We used 3-dimensional surface plots to visualize how these parameters in concert affected our five accuracy metrics and the estimated number of clusters generated by DBSCAN (Figure 17).

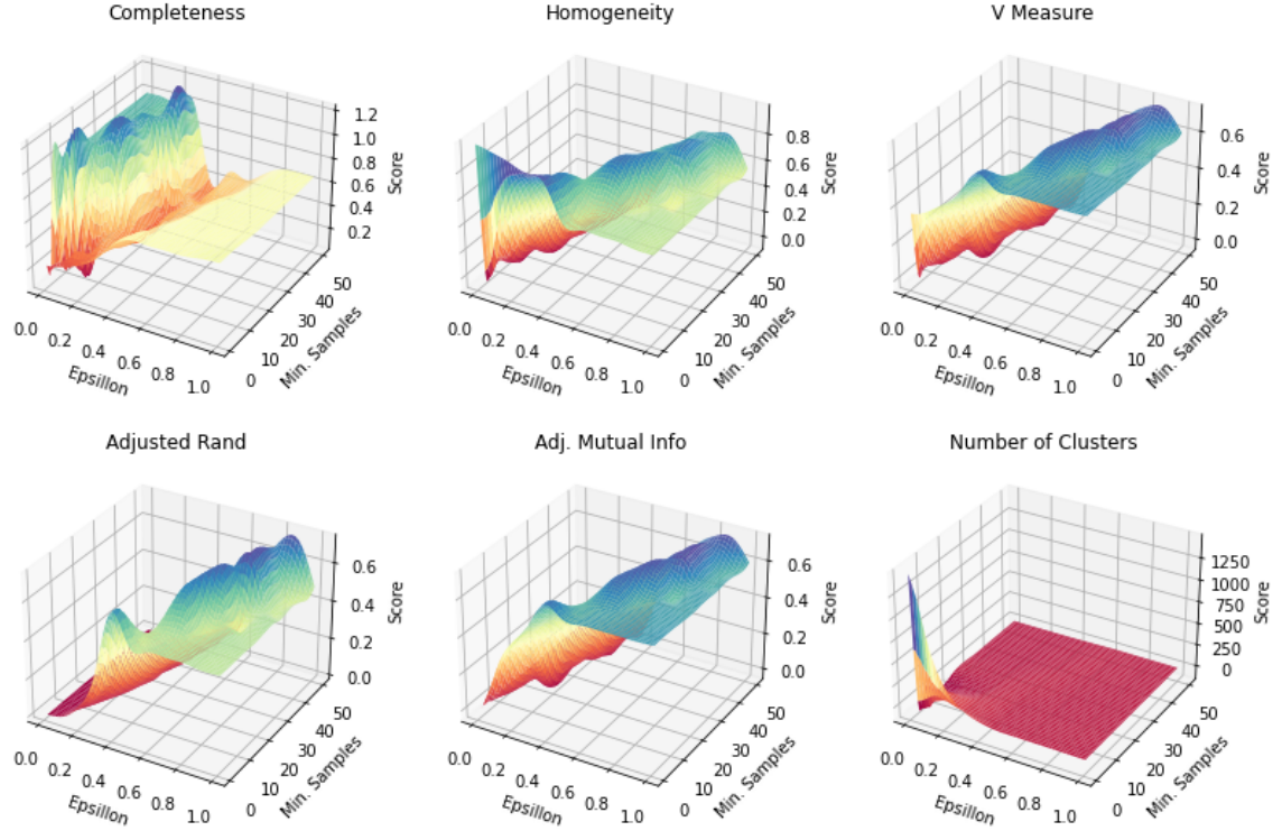


Figure 17: Variation in accuracy metrics corresponding to changes in epsilon and minimum cluster metrics for DBSCAN clustering of UMAP reduced data with Cosine distance metric (Euclidean metric results visually similar)

Avoiding the low values of epsilon and minimum samples that resulted in excessively high numbers of estimated clusters, we found the combination of parameters that maximized the average of the accuracy metrics to be 0.8 for epsilon and 35 for minimum samples. We then compared the results of each of our protocols (Kmeans with Euclidean/Cosine, DBSCAN with Euclidean/Cosine) to find the best reduction and clustering technique for our data. We report the five accuracy metrics for each protocol below:

| <b>Accuracy Metrics for DBSCAN, Metric = Euclidean</b> |                    |
|--|--------------------|
| Homogeneity score                                      | 0.7896799513244787 |
| Completeness score                                     | 0.667782538810624  |
| V_measure score  | 0.7236336939202158 |
| Adjusted Rand Index score                              | 0.7108732847743878 |
| Adjusted Mutual Information Score                      | 0.7221232251101021 |

| <b>Accuracy Metrics for DBSCAN, Metric = Cosine</b> |                    |
|---|--------------------|
| Homogeneity score                                   | 0.7896799513244787 |
| Completeness score                                  | 0.667782538810624  |
| V_measure score                                     | 0.7236336939202158 |
| Adjusted Rand Index score                           | 0.7108732847743878 |
| Adjusted Mutual Information Score                   | 0.7221232251101021 |

| <b>Accuracy Metrics for KMeans, Metric = Euclidean</b> |                    |
|--|--------------------|
| Homogeneity score                                      | 0.7874360944663693 |
| Completeness score                                     | 0.7874929063651126 |
| V_measure score  | 0.7874644993910623 |
| Adjusted Rand Index score                              | 0.8229763601775667 |
| Adjusted Mutual Information Score                      | 0.786746935945491  |

| <b>Accuracy Metrics for KMeans, Metric = Cosine</b> |                    |
|---|--------------------|
| Homogeneity score                                   | 0.7863911146778318 |
| Completeness score                                  | 0.7863053396249526 |
| V_measure score                                     | 0.7863482248123015 |
| Adjusted Rand Index score                           | 0.8206748429226303 |
| Adjusted Mutual Information Score                   | 0.785626959948192  |

The best scoring protocol was Kmeans with Euclidean distance metric and n\_components set to 10. The contingency matrix for this clustering is shown in Figure 18 below

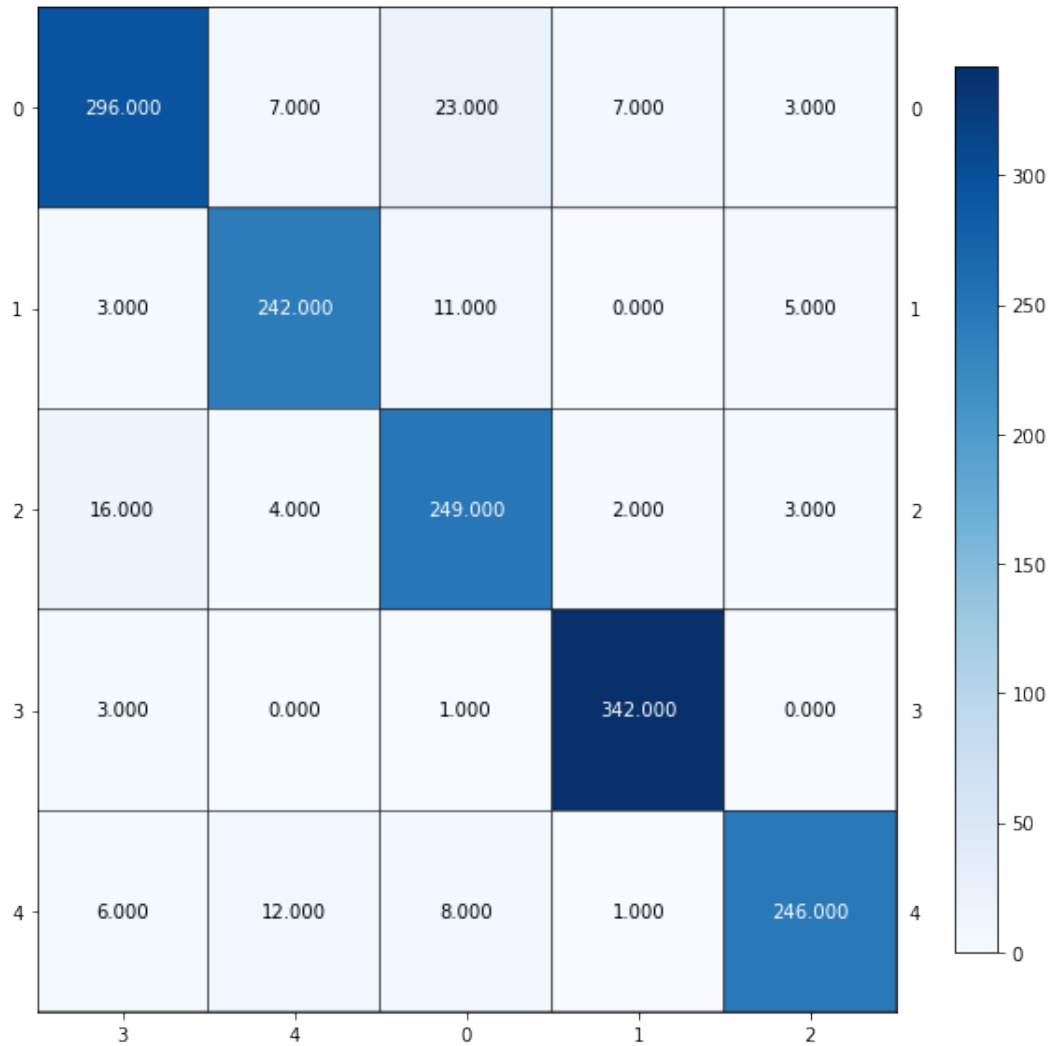


Figure 18: Contingency Matrix for optimal clustering of BBC News data

Our pipeline and parameter selections produce a fairly good clustering through both DBSCAN and KMeans. One factor that may affect the accuracy of our analysis is the relatively small size of the BBC dataset when compared to the 20 Newsgroups dataset. With more data, our clustering would likely become more accurate.

# Project\_2

February 3, 2021

```
[88]: #Q1

import nltk
from nltk import pos_tag
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import CountVectorizer
from string import punctuation
from nltk.corpus import stopwords
from matplotlib import pyplot as plt
from sklearn.datasets import fetch_20newsgroups

categories = ['comp.graphics', 'comp.os.ms-windows.misc',
'comp.sys.ibm.pc.hardware', 'comp.sys.mac.hardware',
'rec.autos', 'rec.motorcycles',
'rec.sport.baseball', 'rec.sport.hockey']

#Vectorizer
analyzer = CountVectorizer().build_analyzer()
tfidf_transformer = TfidfTransformer()

combined_stopwords = set.union(set(stopwords.words('english')),set(punctuation))

def stem_rmv_punc_nolem(doc):
    return (word for word in analyzer(doc) if word not in combined_stopwords and
    ↪not word.isdigit())

#Download dataset
dataset = fetch_20newsgroups(subset='all', categories = categories, shuffle =
    ↪True, random_state = 0, remove=('headers','footers'))

count_vect = CountVectorizer(min_df=3,analyzer=stem_rmv_punc_nolem,
    ↪stop_words='english')

X_counts = count_vect.fit_transform(dataset.data)
X_tfidf = tfidf_transformer.fit_transform(X_counts)
```

```
print("X_tfidf shape: ", X_tfidf.shape)
```

X\_tfidf shape: (7882, 21909)


```
[89]: #Q2
      #Get labels to verify k means accuracy

      from sklearn.cluster import KMeans
      from plotmat import plot_mat
      from sklearn.metrics.cluster import contingency_matrix

      dataset_targets_bin = dataset.target.copy()
      for i in range(len(dataset_targets_bin)):
          if dataset.target[i] in [0,1,2,3]:
              dataset_targets_bin[i] = 0
          else:
              dataset_targets_bin[i] = 1

      kmeans_clf = KMeans(n_clusters=2, max_iter=1500, random_state=0, n_init=50)
      kmeans_clf.fit(X_tfidf)
      km_cont_mat = contingency_matrix(dataset_targets_bin, kmeans_clf.labels_)

      plot_mat(km_cont_mat, size=(5,5))
```



[90]: #Q3

```
from sklearn import metrics
print("Homogeneity score: ", metrics.homogeneity_score(dataset_targets_bin,
    ↳kmeans_clf.labels_))
print("Completeness score: ", metrics.completeness_score(dataset_targets_bin,
    ↳kmeans_clf.labels_))
print("V_measure score: ", metrics.v_measure_score(dataset_targets_bin,
    ↳kmeans_clf.labels_))
print("Adjusted Rand Index score: ", metrics.
    ↳adjusted_rand_score(dataset_targets_bin, kmeans_clf.labels_))
print("Adjusted Mutual Information Score: ", metrics.
    ↳adjusted_mutual_info_score(dataset_targets_bin, kmeans_clf.labels_))
```

Homogeneity score: 0.5600025412974119  
 Completeness score: 0.5716557539454379  
 V\_measure score: 0.5657691483417876  
 Adjusted Rand Index score: 0.6327475635535726

Adjusted Mutual Information Score: 0.5657289846874574

```
[91]: #Q4


from sklearn.decomposition import TruncatedSVD, NMF
import numpy as np

variances = []
r_vals = np.arange(1, 1001)

trunc_svd = TruncatedSVD(n_components=1000, random_state=0)
X_svd = trunc_svd.fit_transform(X_tfidf)

variances = trunc_svd.explained_variance_ratio_

plt.scatter(r_vals, np.cumsum(variances))
plt.xlabel("# of Principle Components")
plt.ylabel("Percent of Variance")
plt.title("Percent of Total Variance vs. Number of Principle Components")
plt.show()
```



```
[92]: #Q5

r_vals_5 = [1,2,3,5,10,20,50,100,300]
metric_scores_svd = []
```



```

metric_scores_nmf = []

for r in r_vals_5:
    nmf_model = NMF(n_components=r, init='random', random_state=0, max_iter=1000)
    nmf_red = nmf_model.fit_transform(X_tfidf)
    svd_red = X_svd[:, :r]


    kmeans_clf.fit(nmf_red)
    metric_scores_nmf.append((metrics.homogeneity_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.completeness_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.v_measure_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.adjusted_rand_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.
↪adjusted_mutual_info_score(dataset_targets_bin, kmeans_clf.labels_)))
    kmeans_clf.fit(svd_red)
    metric_scores_svd.append((metrics.homogeneity_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.completeness_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.v_measure_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.adjusted_rand_score(dataset_targets_bin,
↪kmeans_clf.labels_),
                               metrics.
↪adjusted_mutual_info_score(dataset_targets_bin, kmeans_clf.labels_)))

```


```

[93]: plt.plot(r_vals_5, metric_scores_nmf)
plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
↪ "Adjusted Mutual Info"))
plt.xlabel("r")
plt.ylabel("Score")
plt.title("Scores for Various r using NMF for Dim Reduction")
plt.show()

```



```
[94]: plt.plot(r_vals_5, metric_scores_svd)
plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
           ↪ "Adjusted Mutual Info"))
plt.xlabel("r")
plt.ylabel("Score")
plt.title("Scores for Various r using SVD for Dim Reduction")
plt.show()
```



[95]: #Q7

```
nmf_best_k = 2
svd_best_k = 50

trunc_svd_q3 = TruncatedSVD(n_components=2, random_state=0)

nmf_model = NMF(n_components=nmf_best_k, init='random', random_state=0,
    max_iter=1000)
nmf_red = nmf_model.fit_transform(X_tfidf)
nmf_red_SVD = nmf_red#trunc_svd_q3.fit_transform(nmf_red)
kmeans_clf.fit(nmf_red_SVD)
nmf_labels = kmeans_clf.labels_

svd_red = X_svd[:, :svd_best_k]
svd_red_SVD = trunc_svd_q3.fit_transform(svd_red)
kmeans_clf.fit(svd_red_SVD)
svd_labels = kmeans_clf.labels_

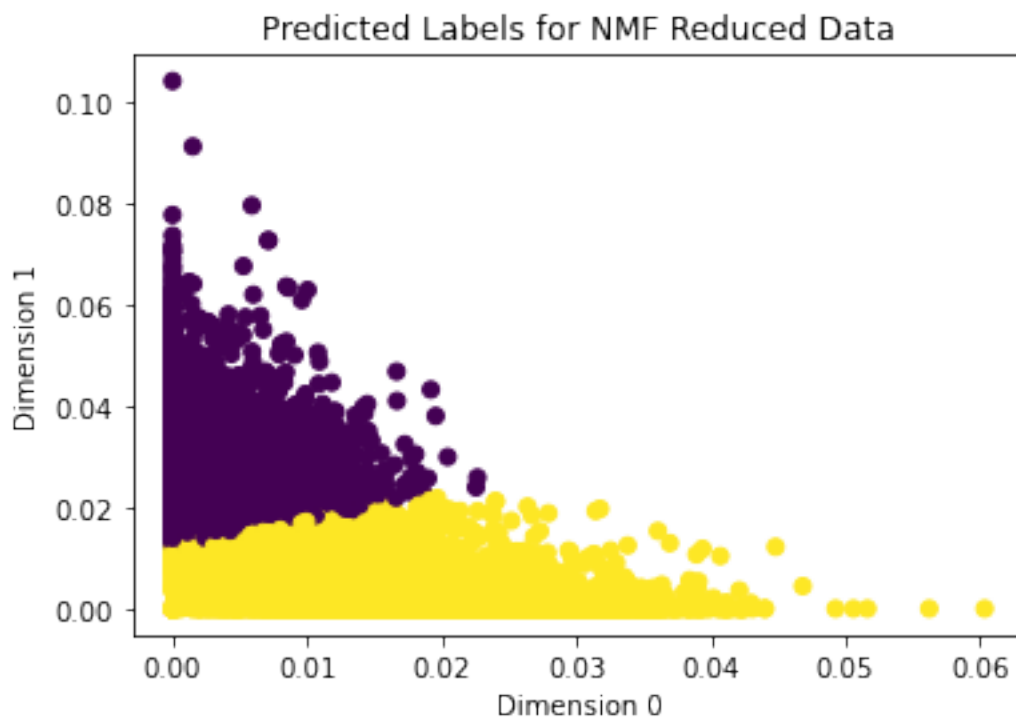
plt.scatter(nmf_red_SVD[:, 0], nmf_red_SVD[:, 1], c=nmf_labels)
plt.title("Predicted Labels for NMF Reduced Data")
plt.xlabel("Dimension 0")
plt.ylabel("Dimension 1")
```



```
plt.show()


plt.scatter(nmf_red_SVD[:,0], nmf_red_SVD[:,1], c=dataset_targets_bin)
plt.title("True Labels for NMF Reduced Data")
plt.xlabel("Dimension 0")
plt.ylabel("Dimension 1")
plt.show()

plt.scatter(svd_red_SVD[:,0], svd_red_SVD[:,1], c=svd_labels)
plt.title("Predicted Labels for SVD Reduced Data")
plt.xlabel("Dimension 0")
plt.ylabel("Dimension 1")
plt.show()

plt.scatter(svd_red_SVD[:,0], svd_red_SVD[:,1], c=dataset_targets_bin)
plt.title("True Labels for SVD Reduced Data")
plt.xlabel("Dimension 0")
plt.ylabel("Dimension 1")
plt.show()
```







```
[96]: #Q9
#get all 20 categories this time
full_dataset = fetch_20newsgroups(subset='all', shuffle = True, random_state = 0,
    remove=('headers','footers'))
full_count_vect = CountVectorizer(min_df=3, analyzer=stem_rmv_punc_nolem,
    stop_words='english')
full_X_counts = full_count_vect.fit_transform(full_dataset.data)
full_X_tfidf = tfidf_transformer.fit_transform(full_X_counts)

#using SVD with r=50 - best results from earlier
full_trunc_svd = TruncatedSVD(n_components=1000, random_state=0)
full_X_svd = full_trunc_svd.fit_transform(full_X_tfidf)
full_svd_red = full_X_svd[:, :50]

#k-means w/20 clusters
full_kmeans_clf = KMeans(n_clusters=20, max_iter=1500, random_state=0, n_init=50)
full_kmeans_clf.fit(full_svd_red)

#metrics!
from scipy.optimize import linear_sum_assignment
from sklearn.metrics import confusion_matrix


cm = contingency_matrix(full_dataset.target, full_kmeans_clf.labels_)
```

```

rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
        size=(15,15), xlabel="Cluster Number", ylabel="Ground Truth Label")

print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
        full_kmeans_clf.labels_))
print("Completeness score: ", metrics.completeness_score(full_dataset.target,
        full_kmeans_clf.labels_))
print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
        full_kmeans_clf.labels_))
print("Adjusted Rand Index score: ", metrics.adjusted_rand_score(full_dataset.
        target, full_kmeans_clf.labels_))
print("Adjusted Mutual Information Score: ",metrics.
        adjusted_mutual_info_score(full_dataset.target, full_kmeans_clf.labels_))

```



Homogeneity score: 0.31630458145750995  
Completeness score: 0.3682552023183464  
V\_measure score: 0.3403086491478565  
Adjusted Rand Index score: 0.10379799750286293  
Adjusted Mutual Information Score: 0.3380068278889245

[97]: #Q10

```
r_vals = [2,5,10,30,50,100]

for i in r_vals:
    nmf_KL = NMF(n_components=i, init='random', random_state=0, max_iter=1000,
    →beta_loss = "kullback-leibler", solver='mu')
    nmf_red_KL = nmf_KL.fit_transform(full_X_tfidf)
    kmeans_clf_KL = KMeans(n_clusters=20, max_iter=1500, random_state=0,
    →n_init=50)
    kmeans_clf_KL.fit(nmf_red_KL)
    print()
    print("r = " + str(i))
    print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
    →kmeans_clf_KL.labels_))
    print("Completeness score: ", metrics.completeness_score(full_dataset.
    →target, kmeans_clf_KL.labels_))
    print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
    →kmeans_clf_KL.labels_))
    print("Adjusted Rand Index score: ", metrics.
    →adjusted_rand_score(full_dataset.target, kmeans_clf_KL.labels_))
    print("Adjusted Mutual Information Score: ",metrics.
    →adjusted_mutual_info_score(full_dataset.target, kmeans_clf_KL.labels_))
```

r = 2  
Homogeneity score: 0.19676527887644293  
Completeness score: 0.2136834606423116  
V\_measure score: 0.2048756966527202  
Adjusted Rand Index score: 0.0604569445452653  
Adjusted Mutual Information Score: 0.20217480540664073

r = 5  
Homogeneity score: 0.35641214262487414  
Completeness score: 0.38144347646289517  
V\_measure score: 0.3685032226345344  
Adjusted Rand Index score: 0.15027690007707672



Adjusted Mutual Information Score: 0.3663744867466823

r = 10

Homogeneity score: 0.37714629232162333

Completeness score: 0.3900505160910544

V\_measure score: 0.3834898799077068

Adjusted Rand Index score: 0.18833101982676798

Adjusted Mutual Information Score: 0.38146535004967785

r = 30

Homogeneity score: 0.36866650441353743

Completeness score: 0.42527782891459665

V\_measure score: 0.39495386265506316

Adjusted Rand Index score: 0.12617792572471453

Adjusted Mutual Information Score: 0.3928569966801935

r = 50

Homogeneity score: 0.30640759897200137

Completeness score: 0.47873972223954575

V\_measure score: 0.3736610566220568

Adjusted Rand Index score: 0.05415577061481132

Adjusted Mutual Information Score: 0.37118024804272814

r = 100

Homogeneity score: 0.1924960903730588

Completeness score: 0.406934904386975

V\_measure score: 0.26135911828245

Adjusted Rand Index score: 0.028607969901381668

Adjusted Mutual Information Score: 0.25806962570480285

```
[153]: #Q11

import umap

n_vals = [2,5,10,50,100,200,300,500,700,1000]
metric_scores_euc = []
metric_scores_cos = []

for i in n_vals:
    print("working on reducing for n = " + str(i))
    reducer_euc = umap.UMAP(n_components=i, init='random', metric="euclidean")
    embedding_euc = reducer_euc.fit_transform(full_X_tfidf)
    reducer_cos = umap.UMAP(n_components=i, init='random', metric="cosine")
    embedding_cos = reducer_cos.fit_transform(full_X_tfidf)

    print("working on kmeans for n = " + str(i))
```

```

    kmeans_clf_euc = KMeans(n_clusters=20, max_iter=1500, random_state=0,
→n_init=50)
    kmeans_clf_euc.fit(embedding_euc)
    kmeans_clf_cos = KMeans(n_clusters=20, max_iter=1500, random_state=0,
→n_init=50)
    kmeans_clf_cos.fit(embedding_cos)

    metric_scores_euc.append((metrics.homogeneity_score(full_dataset.target,
→kmeans_clf_euc.labels_),
                                metrics.completeness_score(full_dataset.target,
→kmeans_clf_euc.labels_),
                                metrics.v_measure_score(full_dataset.target,
→kmeans_clf_euc.labels_),
                                metrics.adjusted_rand_score(full_dataset.target,
→kmeans_clf_euc.labels_),
                                metrics.adjusted_mutual_info_score(full_dataset.
→target, kmeans_clf_euc.labels_)))

    metric_scores_cos.append((metrics.homogeneity_score(full_dataset.target,
→kmeans_clf_cos.labels_),
                                metrics.completeness_score(full_dataset.target,
→kmeans_clf_cos.labels_),
                                metrics.v_measure_score(full_dataset.target,
→kmeans_clf_cos.labels_),
                                metrics.adjusted_rand_score(full_dataset.target,
→kmeans_clf_cos.labels_),
                                metrics.adjusted_mutual_info_score(full_dataset.
→target, kmeans_clf_cos.labels_)))

plt.plot(n_vals, metric_scores_euc)
plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
→"Adjusted Mutual Info"))
plt.xlabel("n_components")
plt.ylabel("Score")
plt.title("Scores for Various n_components Using UMAP w/Euclidean Metric")
plt.show()

plt.plot(n_vals, metric_scores_cos)
plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
→"Adjusted Mutual Info"))
plt.xlabel("n_components")
plt.ylabel("Score")
plt.title("Scores for Various n_components Using UMAP w/Cosine Metric")
plt.show()

```

working on reducing for n = 2

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-153-abf5f6750690> in <module>()
    11     print("working on reducing for n = " + str(i))
    12     reducer_euc = umap.UMAP(n_components=i, init='random',
    ↪metric="euclidean")
--> 13     embedding_euc = reducer_euc.fit_transform(full_X_tfidf)
    14     reducer_cos = umap.UMAP(n_components=i, init='random', metric="cosine")
    15     embedding_cos = reducer_cos.fit_transform(full_X_tfidf)

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\umap\umap_.py in
    ↪fit_transform(self, X, y)
    2012         Embedding of the training data in low-dimensional space.
    2013         """
-> 2014         self.fit(X, y)
    2015         return self.embedding_
    2016

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\umap\umap_.py in
    ↪fit(self, X, y)
    1799         self.low_memory,
    1800         use_pynndescent=True,
-> 1801         verbose=self.verbose,
    1802     )
    1803

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\umap\umap_.py in
    ↪nearest_neighbors(X, n_neighbors, metric, metric_kwds, angular, random_state,
    ↪low_memory, use_pynndescent, verbose)
    298         max_candidates=60,
    299         low_memory=low_memory,
--> 300         verbose=verbose,
    301     )
    302     knn_indices, knn_dists = nnd.neighbor_graph

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\pynndescent\pynndescent_.py in
    ↪__init__(self, data, metric, metric_kwds, n_neighbors, n_trees,
    ↪leaf_size, pruning_degree_multiplier, diversify_prob, n_search_trees,
    ↪tree_init, init_graph, random_state, low_memory, max_candidates, n_iters,
    ↪delta, n_jobs, compressed, verbose)
    789         current_random_state,
    790         self.n_jobs,
--> 791         self._angular_trees,
    792     )
    793     leaf_array = rptree_leaf_array(self._rp_forest)
```

```

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\pynndescent\rp_trees.py
→in make_forest(data, n_neighbors, n_trees, leaf_size, rng_state, random_state,
→n_jobs, angular)
    994             angular,
    995         )
--> 996         for i in range(n_trees)
    997     )
    998     else:

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\joblib\parallel.py in
→__call__(self, iterable)
    1059
    1060         with self._backend.retrieval_context():
-> 1061             self.retrieve()
    1062         # Make sure that we get a last message telling us we are done
    1063         elapsed_time = time.time() - self._start_time

C:\Users\loicm\anaconda3\envs\ece219_2\lib\site-packages\joblib\parallel.py in
→retrieve(self)
    938         try:
    939             if getattr(self._backend, 'supports_timeout', False):
--> 940                 self._output.extend(job.get(timeout=self.timeout))
    941             else:
    942                 self._output.extend(job.get())

C:\Users\loicm\anaconda3\envs\ece219_2\lib\multiprocessing\pool.py in get(self,
→timeout)
    649
    650     def get(self, timeout=None):
--> 651         self.wait(timeout)
    652         if not self.ready():
    653             raise TimeoutError

C:\Users\loicm\anaconda3\envs\ece219_2\lib\multiprocessing\pool.py in wait(self,
→timeout)
    646
    647     def wait(self, timeout=None):
--> 648         self._event.wait(timeout)
    649
    650     def get(self, timeout=None):

C:\Users\loicm\anaconda3\envs\ece219_2\lib\threading.py in wait(self, timeout)
    550         signaled = self._flag
    551         if not signaled:
--> 552             signaled = self._cond.wait(timeout)
    553         return signaled
    554

```

```

C:\Users\loicm\anaconda3\envs\ece219_2\lib\threading.py in wait(self, timeout)
    294         try:      # restore state no matter what (e.g., KeyboardInterrupt)
    295             if timeout is None:
--> 296                 waiter.acquire()
    297                 gotit = True
    298             else:

```

KeyboardInterrupt:

```

[143]: #Q12
#Best n_components = 5 for both euclidean and cosine metrics

reducer_euc = umap.UMAP(n_components=5, init='random', metric="euclidean",
    ↳random_state=42)
embedding_euc = reducer_euc.fit_transform(full_X_tfidf)
reducer_cos = umap.UMAP(n_components=5, init='random', metric="cosine",
    ↳random_state=42)
embedding_cos = reducer_cos.fit_transform(full_X_tfidf)

kmeans_clf_euc = KMeans(n_clusters=20, max_iter=1500, random_state=0, n_init=50)
kmeans_clf_euc.fit(embedding_euc)
kmeans_clf_cos = KMeans(n_clusters=20, max_iter=1500, random_state=0, n_init=50)
kmeans_clf_cos.fit(embedding_cos)

print("Performance of UMAP reduced data for n_components=5")
print()
print("Euclidean metric")
print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
    ↳kmeans_clf_euc.labels_))
print("Completeness score: ", metrics.completeness_score(full_dataset.target,
    ↳kmeans_clf_euc.labels_))
print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
    ↳kmeans_clf_euc.labels_))
print("Adjusted Rand Index score: ", metrics.adjusted_rand_score(full_dataset.
    ↳target, kmeans_clf_euc.labels_))
print("Adjusted Mutual Information Score: ",metrics.
    ↳adjusted_mutual_info_score(full_dataset.target, kmeans_clf_euc.labels_))

print()
print("Cosine metric")
print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
    ↳kmeans_clf_cos.labels_))
print("Completeness score: ", metrics.completeness_score(full_dataset.target,
    ↳kmeans_clf_cos.labels_))
print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
    ↳kmeans_clf_cos.labels_))

```

```

print("Adjusted Rand Index score: ", metrics.adjusted_rand_score(full_dataset.
    →target, kmeans_clf_cos.labels_))
print("Adjusted Mutual Information Score: ", metrics.
    →adjusted_mutual_info_score(full_dataset.target, kmeans_clf_cos.labels_))

cm = contingency_matrix(full_dataset.target, kmeans_clf_euc.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
    →size=(15,15), xlabel="Cluster Number", ylabel="Ground Truth Label")

cm = contingency_matrix(full_dataset.target, kmeans_clf_cos.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
    →size=(15,15), xlabel="Cluster Number", ylabel="Ground Truth Label")

```

Performance of UMAP reduced data for n\_components=5

Euclidean metric

Homogeneity score: 0.008428056566878686

Completeness score: 0.008490790660743977

V\_measure score: 0.008459307306639803

Adjusted Rand Index score: 0.0014762057832842827

Adjusted Mutual Information Score: 0.00524781170211466

Cosine metric


Homogeneity score: 0.5347926404764849


Completeness score: 0.5549487009354954

V\_measure score: 0.5446842655666289

Adjusted Rand Index score: 0.4167210932677909

Adjusted Mutual Information Score: 0.5431756896373161





```
[142]: #Q13
#using same umap parameters as above

from sklearn.cluster import AgglomerativeClustering

umap_best_n = 5
umap_best_metric = 'cosine'

umap_reducer = umap.UMAP(n_components=umap_best_n, init='random',
    ↪metric=umap_best_metric, random_state=42)

red_umap = umap_reducer.fit_transform(full_X_tfidf)
```



```

agg_clust_ward = AgglomerativeClustering(n_clusters=20, linkage='ward')
agg_clust_single = AgglomerativeClustering(n_clusters=20, linkage='single')

agg_clust_ward.fit(red_umap)
agg_clust_single.fit(red_umap)

print("Ward Linkage Scores")
print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
    ↳agg_clust_ward.labels_))
print("Completeness score: ", metrics.completeness_score(full_dataset.target,
    ↳agg_clust_ward.labels_))
print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
    ↳agg_clust_ward.labels_))
print("Adjusted Rand Index score: ", metrics.adjusted_rand_score(full_dataset.
    ↳target, agg_clust_ward.labels_))
print("Adjusted Mutual Information Score: ",metrics.
    ↳adjusted_mutual_info_score(full_dataset.target, agg_clust_ward.labels_))

print()
print("Single Linkage Scores")
print("Homogeneity score: ", metrics.homogeneity_score(full_dataset.target,
    ↳agg_clust_single.labels_))
print("Completeness score: ", metrics.completeness_score(full_dataset.target,
    ↳agg_clust_single.labels_))
print("V_measure score: ", metrics.v_measure_score(full_dataset.target,
    ↳agg_clust_single.labels_))
print("Adjusted Rand Index score: ", metrics.adjusted_rand_score(full_dataset.
    ↳target, agg_clust_single.labels_))
print("Adjusted Mutual Information Score: ",metrics.
    ↳adjusted_mutual_info_score(full_dataset.target, agg_clust_single.labels_))

```

#### Ward Linkage Scores

Homogeneity score: 0.524473300498607  
 Completeness score: 0.5479912976714623  
 V\_measure score: 0.53597443687123  
 Adjusted Rand Index score: 0.3940501791969448  
 Adjusted Mutual Information Score: 0.5344316766896055

#### Single Linkage Scores

Homogeneity score: 0.022917457355135236  
 Completeness score: 0.3464156710750703  
 V\_measure score: 0.04299081646294147  
 Adjusted Rand Index score: 0.00046761208355549543  
 Adjusted Mutual Information Score: 0.03748616963750802

[140]: *#Q14 with more hyperparams*

```
eps_vals = [.01,.05,.1,.3,.5,.7,.9,1]
metric_types = [ 'euclidean', 'manhattan', 'cosine']
min_samples = [10, 20, 50, 70, 100]

db_max_score = -100
hdb_max_score = -100

for eps in eps_vals:
    for metric_t in metric_types:
        for min_samp in min_samples:
            print("Current Iteration: eps={:f}, metric={:s}, min_samples={:d}".
→format(eps, metric_t, min_samp))
            DBSCAN_clust = DBSCAN(min_samples=min_samp, metric=metric_t, eps=eps)

            #HDB does not support cosine metric
            if(metric_t != 'cosine'):
                HDBSCAN_clust = HDBSCAN(min_cluster_size=100,
→min_samples=min_samp, cluster_selection_epsilon=eps)

                DBSCAN_clust.fit(red_umap)
                HDBSCAN_clust.fit(red_umap)

                homog_score = metrics.homogeneity_score(full_dataset.target,
→DBSCAN_clust.labels_)
                comp_score = metrics.completeness_score(full_dataset.target,
→DBSCAN_clust.labels_)
                v_score = metrics.v_measure_score(full_dataset.target, DBSCAN_clust.
→labels_)
                rand_score = metrics.adjusted_rand_score(full_dataset.target,
→DBSCAN_clust.labels_)
                adj_mut_score = metrics.adjusted_mutual_info_score(full_dataset.
→target, DBSCAN_clust.labels_)

                score_avg = (homog_score + comp_score + v_score + rand_score +
→adj_mut_score)/5

                labels = DBSCAN_clust.labels_
                n_clust = len(set(labels)) - (1 if -1 in labels else 0)

                if score_avg > db_max_score:
                    db_max_score = score_avg
                    db_max_params = (eps, metric_t, min_samp)
                    db_max_metrics = (homog_score, comp_score, v_score, rand_score,
→adj_mut_score)
```

```

        homog_score = metrics.homogeneity_score(full_dataset.target,
→HDBSCAN_clust.labels_)
        comp_score = metrics.completeness_score(full_dataset.target,
→HDBSCAN_clust.labels_)
        v_score = metrics.v_measure_score(full_dataset.target, HDBSCAN_clust.
→labels_)
        rand_score = metrics.adjusted_rand_score(full_dataset.target,
→HDBSCAN_clust.labels_)
        adj_mut_score = metrics.adjusted_mutual_info_score(full_dataset.
→target, HDBSCAN_clust.labels_)

        score_avg = (homog_score + comp_score + v_score + rand_score +
→adj_mut_score)/5

        labels = HDBSCAN_clust.labels_
        n_clust = len(set(labels)) - (1 if -1 in labels else 0)

        if metric_t != 'cosine' and score_avg > hdb_max_score:
            hdb_max_score = score_avg
            hdb_max_params = (eps, metric_t, min_samp)
            hdb_max_metrics = (homog_score, comp_score, v_score, rand_score,
→adj_mut_score)

print()
print("DBSCAN Best Parameters: eps={:f}, metric={:s}, min_samples={:d}".
→format(db_max_params[0], db_max_params[1], db_max_params[2]))
print()
print("Best DBSCAN Scores")
print("Homogeneity score: ", db_max_metrics[0])
print("Completeness score: ", db_max_metrics[1])
print("V_measure score: ", db_max_metrics[2])
print("Adjusted Rand Index score: ", db_max_metrics[3])
print("Adjusted Mutual Information Score: ", db_max_metrics[4])

print()
print("HDBSCAN Best Parameters: eps={:f}, metric={:s}, min_samples={:d}".
→format(hdb_max_params[0], hdb_max_params[1], hdb_max_params[2]))
print()
print("Best HDBSCAN Scores")
print("Homogeneity score: ", hdb_max_metrics[0])
print("Completeness score: ", hdb_max_metrics[1])
print("V_measure score: ", hdb_max_metrics[2])
print("Adjusted Rand Index score: ", hdb_max_metrics[3])
print("Adjusted Mutual Information Score: ", hdb_max_metrics[4])

```



[illegible]

```

Current Iteration: eps=0.900000, metric=manhattan, min_samples=20
Current Iteration: eps=0.900000, metric=manhattan, min_samples=50
Current Iteration: eps=0.900000, metric=manhattan, min_samples=70
Current Iteration: eps=0.900000, metric=manhattan, min_samples=100
Current Iteration: eps=0.900000, metric=cosine, min_samples=10
Current Iteration: eps=0.900000, metric=cosine, min_samples=20
Current Iteration: eps=0.900000, metric=cosine, min_samples=50
Current Iteration: eps=0.900000, metric=cosine, min_samples=70
Current Iteration: eps=0.900000, metric=cosine, min_samples=100
Current Iteration: eps=1.000000, metric=euclidean, min_samples=10
Current Iteration: eps=1.000000, metric=euclidean, min_samples=20
Current Iteration: eps=1.000000, metric=euclidean, min_samples=50
Current Iteration: eps=1.000000, metric=euclidean, min_samples=70
Current Iteration: eps=1.000000, metric=euclidean, min_samples=100
Current Iteration: eps=1.000000, metric=manhattan, min_samples=10
Current Iteration: eps=1.000000, metric=manhattan, min_samples=20
Current Iteration: eps=1.000000, metric=manhattan, min_samples=50
Current Iteration: eps=1.000000, metric=manhattan, min_samples=70
Current Iteration: eps=1.000000, metric=manhattan, min_samples=100
Current Iteration: eps=1.000000, metric=cosine, min_samples=10
Current Iteration: eps=1.000000, metric=cosine, min_samples=20
Current Iteration: eps=1.000000, metric=cosine, min_samples=50
Current Iteration: eps=1.000000, metric=cosine, min_samples=70
Current Iteration: eps=1.000000, metric=cosine, min_samples=100

```

DBSCAN Best Parameters: eps=0.500000, metric=manhattan, min\_samples=10

#### Best DBSCAN Scores

```

Homogeneity score: 0.5407607595568981
Completeness score: 0.45471096739764455
V_measure score: 0.49401673890037173
Adjusted Rand Index score: 0.2967260971944622
Adjusted Mutual Information Score: 0.47920988791262803

```

HDBSCAN Best Parameters: eps=0.300000, metric=euclidean, min\_samples=10

#### Best HDBSCAN Scores

```

Homogeneity score: 0.43889528158878366
Completeness score: 0.5288819318072416
V_measure score: 0.47970500064413385
Adjusted Rand Index score: 0.2565568616115164
Adjusted Mutual Information Score: 0.4780526456871693

```

[141]: # Q15

```

DBSCAN_clust = DBSCAN(min_samples=10, eps=0.5, metric='manhattan')
DBSCAN_clust.fit(red_umap)

```


```

cm = contingency_matrix(full_dataset.target, DBSCAN_clust.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
        size=(15,15), xlabel="Cluster Number", ylabel="Ground Truth Label")

#when label is -1, DBSCAN thinks the point is noise - so it might estimate fewer
clusters than we actually have!

labels = DBSCAN_clust.labels_
n_clust = len(set(labels)) - (1 if -1 in labels else 0)
print("Estimated number of clusters: ", n_clust)

```



Estimated number of clusters: 168

[126]: #Q16

```
#get data and vectorize

#Using the larger training dataset
import pandas
from sklearn.cluster import DBSCAN
from sklearn.metrics.cluster import contingency_matrix
from scipy.optimize import linear_sum_assignment
from plotmat import plot_mat
import numpy as np
import umap.umap_ as umap

bbc = pandas.read_csv('BBC_News_Train.csv')
categories = (bbc.Category.unique()).tolist()
targets = []
for i in bbc.Category:
    for j in range(0,len(categories)):
        if i == categories[j]:
            targets.append(j)
bbc['target'] = targets

X_counts = count_vect.fit_transform(bbc.Text)
bbc_X_tfidf = tfidf_transformer.fit_transform(X_counts)
print("bbc_X_tfidf shape: ", bbc_X_tfidf.shape)
```

bbc\_X\_tfidf shape: (1490, 10197)

[150]: #Find best number of components and metric for UMAP reduction, using KMeans  
→scores as evaluation

```
# Find best N and distance metric for KMeans

n_vals = [2,5,10,50,100,1000]
metric_scores_euc = []
metric_scores_cos = []

for i in n_vals:
    print("working on reducing for n = " + str(i))
    reducer_euc = umap.UMAP(n_components=i, init='random', metric="euclidean",
→random_state=0)
    embedding_euc = reducer_euc.fit_transform(bbc_X_tfidf)
    reducer_cos = umap.UMAP(n_components=i, init='random', metric="cosine",
→random_state=0)
    embedding_cos = reducer_cos.fit_transform(bbc_X_tfidf)
```



```

col_mean = np.nanmean(embedding_cos, axis=0)
inds = np.where(np.isnan(embedding_cos))
embedding_cos[inds] = np.take(col_mean, inds[1])

print("working on kmeans for n = " + str(i))
kmeans_clf_euc = KMeans(n_clusters=5, max_iter=1500, random_state=0,
→n_init=50)
kmeans_clf_euc.fit(embedding_euc)
kmeans_clf_cos = KMeans(n_clusters=5, max_iter=1500, random_state=0,
→n_init=50)
kmeans_clf_cos.fit(embedding_cos)

metric_scores_euc.append((metrics.homogeneity_score(bbc.target,
→kmeans_clf_euc.labels_),
                        metrics.completeness_score(bbc.target,
→kmeans_clf_euc.labels_),
                        metrics.v_measure_score(bbc.target, kmeans_clf_euc.
→labels_),
                        metrics.adjusted_rand_score(bbc.target,
→kmeans_clf_euc.labels_),
                        metrics.adjusted_mutual_info_score(bbc.target,
→kmeans_clf_euc.labels_)))

metric_scores_cos.append((metrics.homogeneity_score(bbc.target,
→kmeans_clf_cos.labels_),
                        metrics.completeness_score(bbc.target,
→kmeans_clf_cos.labels_),
                        metrics.v_measure_score(bbc.target, kmeans_clf_cos.
→labels_),
                        metrics.adjusted_rand_score(bbc.target,
→kmeans_clf_cos.labels_),
                        metrics.adjusted_mutual_info_score(bbc.target,
→kmeans_clf_cos.labels_)))

# Plot Metrics against n_components
plt.plot(n_vals, metric_scores_euc)
plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
→"Adjusted Mutual Info"))
plt.xlabel("n_components")
plt.ylabel("Score")
plt.title("Scores for Various n_components Using UMAP w/Euclidean Metric")
plt.show()

plt.plot(n_vals, metric_scores_cos)

```

```

plt.legend(("Homogeneity", "Completeness", "V_measure", "Adjusted Rand Score",
    ↳ "Adjusted Mutual Info"))
plt.xlabel("n_components")
plt.ylabel("Score")
plt.title("Scores for Various n_components Using UMAP w/Cosine Metric")
plt.show()

# Find best values and print metrics

avs_Kmeans_cos = []
for i in range(len(metric_scores_cos)):
    a =
    ↳ ((metric_scores_cos[i][0]+metric_scores_cos[i][1]+metric_scores_cos[i][2]+metric_scores_cos[i]
    ↳ 5)
    avs_Kmeans_cos.append(a)

avs_Kmeans_euc = []
for i in range(len(metric_scores_euc)):
    a =
    ↳ ((metric_scores_euc[i][0]+metric_scores_euc[i][1]+metric_scores_euc[i][2]+metric_scores_euc[i]
    ↳ 5)
    avs_Kmeans_euc.append(a)

print('Best Kmeans Cos: ', max(avs_Kmeans_cos))
print('Best Kmeans Euc: ', max(avs_Kmeans_euc))

c = avs_Kmeans_euc.index(max(avs_Kmeans_euc))
b = avs_Kmeans_cos.index(max(avs_Kmeans_cos))

print()
print("Best KMeans EUC Scores")
print("Homogeneity score: ", metric_scores_euc[c][0])
print("Completeness score: ", metric_scores_euc[c][1])
print("V_measure score: ", metric_scores_euc[c][2])
print("Adjusted Rand Index score: ", metric_scores_euc[c][3])
print("Adjusted Mutual Information Score: ", metric_scores_euc[c][4])
print("Best N_Components: ", n_vals[c])

umap_reducer = umap.UMAP(n_components = n_vals[c], init = 'random', metric =
    ↳ 'euclidean', random_state=0)
bbc_umap = umap_reducer.fit_transform(bbc_X_tfidf)

#k-means w/5 clusters
bbc_kmeans_clf = KMeans(n_clusters=5, max_iter=1500, random_state=0, n_init=50)
bbc_kmeans_clf.fit(bbc_umap)

#metrics!

```

```

from scipy.optimize import linear_sum_assignment
from sklearn.metrics import confusion_matrix

cm = contingency_matrix(bbc.target, bbc_kmeans_clf.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
    ↪size=(8,8))

print()
print("Best KMeans COS Scores")
print("Homogeneity score: ", metric_scores_cos[b][0])
print("Completeness score: ", metric_scores_cos[b][1])
print("V_measure score: ", metric_scores_cos[b][2])
print("Adjusted Rand Index score: ", metric_scores_cos[b][3])
print("Adjusted Mutual Information Score: ", metric_scores_cos[b][4])
print("Best N_Components: ", n_vals[b])

umap_reducer = umap.UMAP(n_components = n_vals[c], init = 'random', metric =
    ↪'cosine', random_state=0)
bbc_umap = umap_reducer.fit_transform(bbc_X_tfidf)

#k-means w/5 clusters
bbc_kmeans_clf = KMeans(n_clusters=5, max_iter=1500, random_state=0, n_init=50)
bbc_kmeans_clf.fit(bbc_umap)

#metrics!
from scipy.optimize import linear_sum_assignment
from sklearn.metrics import confusion_matrix

cm = contingency_matrix(bbc.target, bbc_kmeans_clf.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,
    ↪size=(8,8))

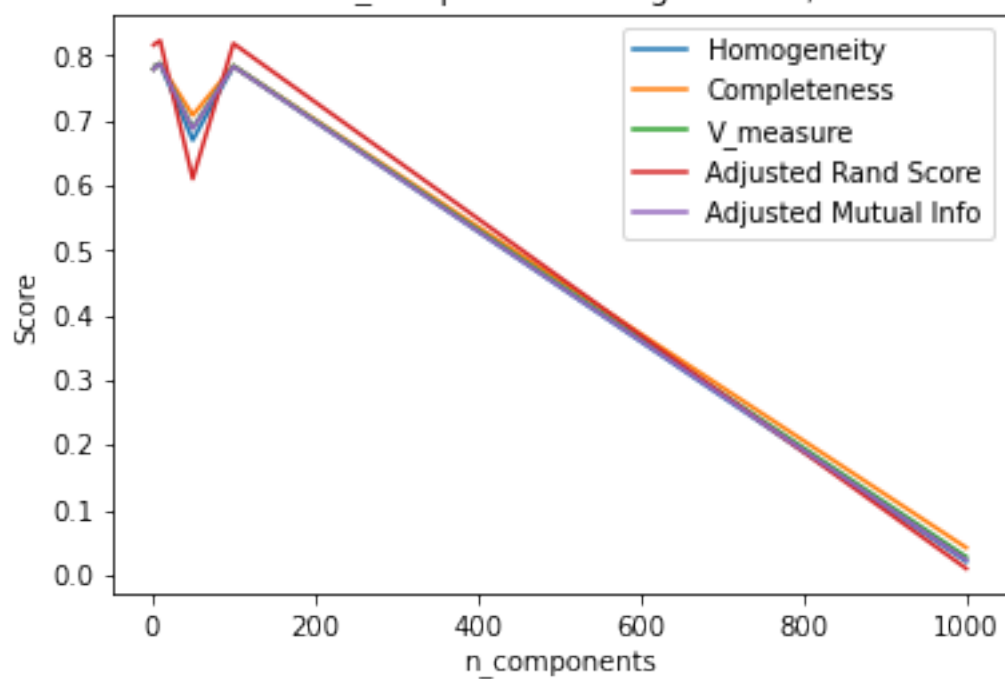
```

```

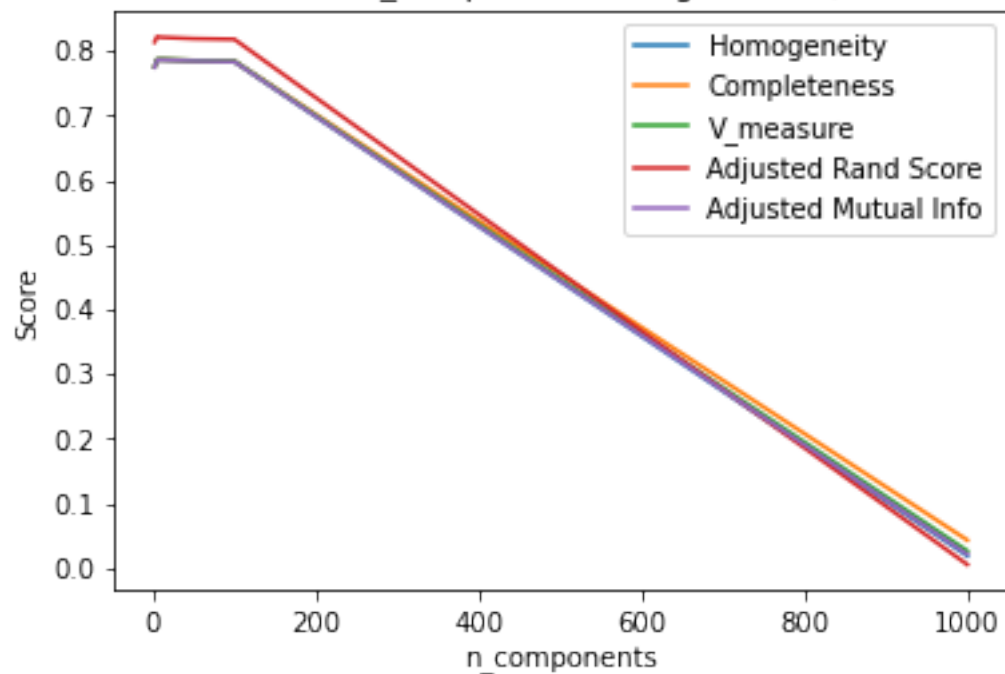
working on reducing for n = 2
working on kmeans for n = 2
working on reducing for n = 5
working on kmeans for n = 5
working on reducing for n = 10
working on kmeans for n = 10
working on reducing for n = 50
working on kmeans for n = 50
working on reducing for n = 100
working on kmeans for n = 100
working on reducing for n = 1000
working on kmeans for n = 1000

```

Scores for Various n\_components Using UMAP w/Euclidean Metric



Scores for Various n\_components Using UMAP w/Cosine Metric



Best Kmeans Cos: 0.7930692963971817

Best Kmeans Euc: 0.7944233592691203

Best KMeans EUC Scores

Homogeneity score: 0.7874360944663693

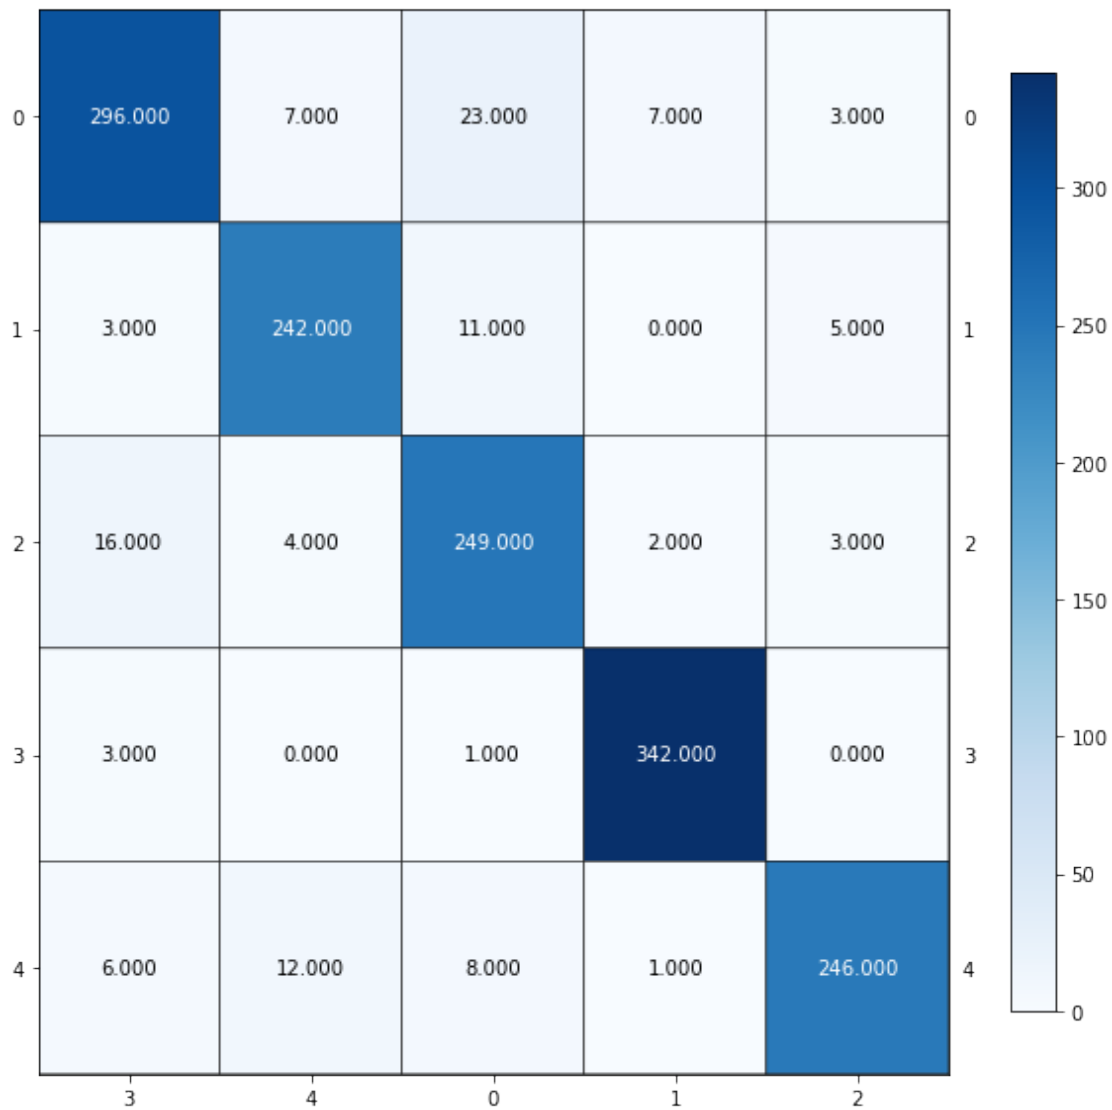
Completeness score: 0.7874929063651126

V\_measure score: 0.7874644993910623

Adjusted Rand Index score: 0.8229763601775667

Adjusted Mutual Information Score: 0.786746935945491

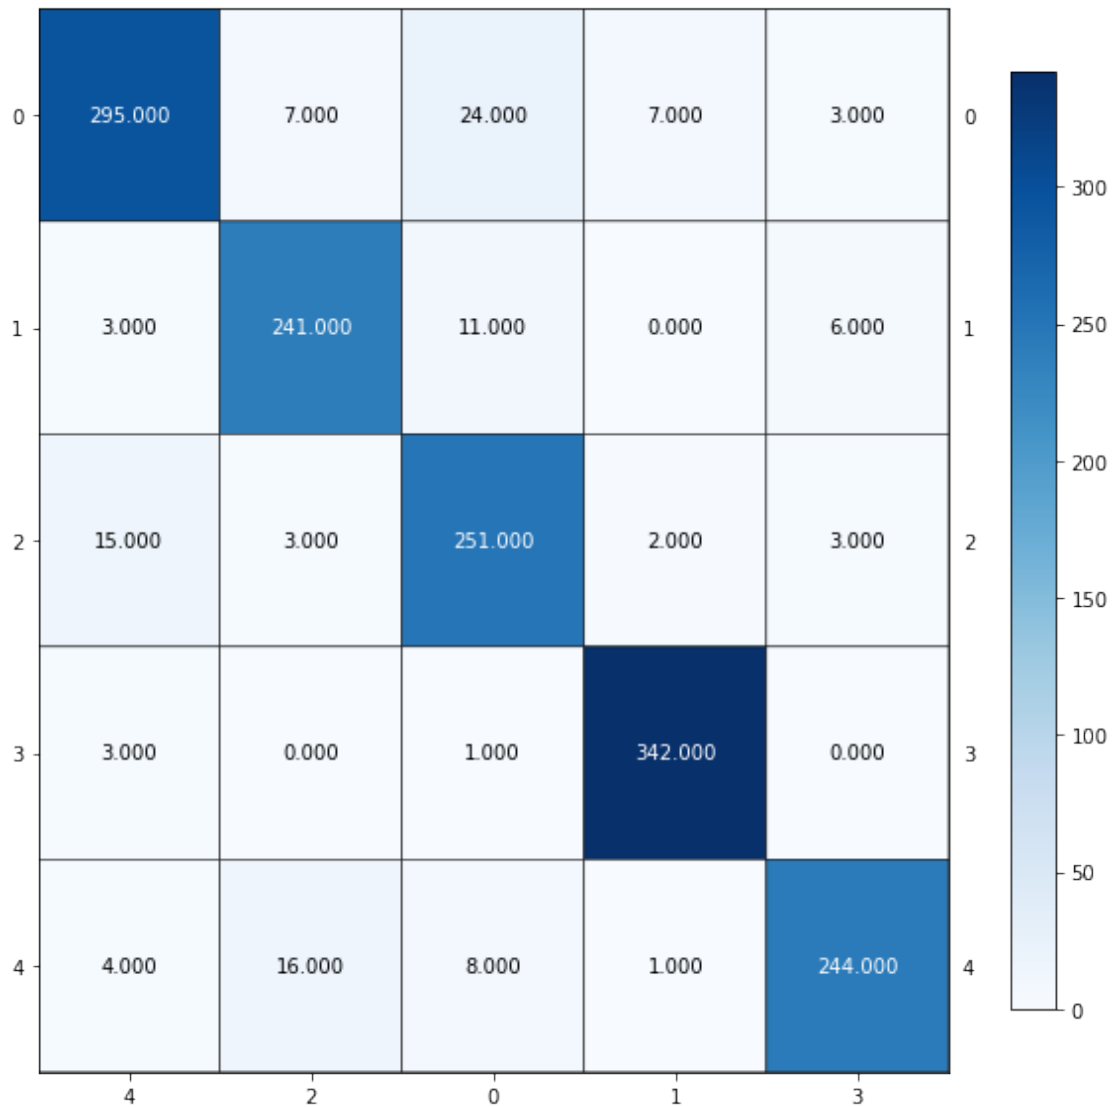
Best N\_Components: 10



Best KMeans COS Scores

Homogeneity score: 0.7863911146778318

Completeness score: 0.7863053396249526  
 V\_measure score: 0.7863482248123015  
 Adjusted Rand Index score: 0.8206748429226303  
 Adjusted Mutual Information Score: 0.785626959948192  
 Best N\_Components: 10



```

[151]: # Find best parameter combo in DBSCAN with COSINE metric
# Generate 3D surface plots for each metric against epsilon and min samples
→ values

umap_reducer = umap.UMAP(n_components = 5, init = 'random', metric = 'cosine',
→ random_state=0)
bbc_umap_COS = umap_reducer.fit_transform(bbc_X_tfidf)
  
```

```

eps_vals = [.01,.05,.1,.2,.3,.4,.5,.6,.7,.8,.9,1]
samples = [1,3,5,10,20,30,35,40,50]
metric_scores_DBSCAN_COS = []

for eps in eps_vals:
    for sam in samples:
        DBSCAN_clust = DBSCAN(min_samples=sam, eps=eps)

        DBSCAN_clust.fit(bbc_umap)

        labels = DBSCAN_clust.labels_
        n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)

        metric_scores_DBSCAN_COS.append([metrics.homogeneity_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.completeness_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.v_measure_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.adjusted_rand_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.adjusted_mutual_info_score(bbc.
→target, DBSCAN_clust.labels_),
                                         n_clusters_,eps,sam])

# Visualize parameter effects on metrics
from mpl_toolkits.mplot3d import Axes3D
from scipy.interpolate import griddata
metricdata = np.array(metric_scores_DBSCAN_COS)
metric_scores = ['Homogeneity', 'Completeness', 'V Measure', 'Adjusted Rand',
→'Adj. Mutual Info', 'Number of Clusters']

for i in range(0,6):
    plotdata = metricdata[:, [i,6,7]]
    z, x, y = zip(*plotdata)
    z = list(map(float, z))
    grid_x, grid_y = np.mgrid[min(x):max(x):100j, min(y):max(y):100j]
    grid_z = griddata((x, y), z, (grid_x, grid_y), method='cubic')

    fig = plt.figure()
    ax = fig.gca(projection='3d')
    ax.plot_surface(grid_x, grid_y, grid_z, cmap=plt.cm.Spectral)
    ax.set_xlabel('Epsilon')
    ax.set_ylabel('Min. Samples')

```

```

ax.set_zlabel('Score')
ax.set_title(str(metric_scores[i]))
plt.show()

# Find best parameter combination for epsilon and minimum samples


avs = []
for i in range(len(metricdata)):
    a = (metricdata[i][0]+metricdata[i][1]+metricdata[i][2]+metricdata[i][3]+metricdata[i][4])/5
    avs.append(a)

b = avs.index(max(avs))
print('Best Epsilon = ', metricdata[b][6])
print('Best N_Samples = ', metricdata[b][7])
print()
cm = contingency_matrix(bbc.target, DBSCAN_clust.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows, size=(5,5))
print("Estimated Number of Clusters: ", metricdata[b][5])
print()
print("Best DBSCAN Scores for Cosine Distance Metric")
print("Homogeneity score: ", metricdata[b][0])
print("Completeness score: ", metricdata[b][1])
print("V_measure score: ", metricdata[b][2])
print("Adjusted Rand Index score: ", metricdata[b][3])
print("Adjusted Mutual Information Score: ", metricdata[b][4])


```




Homogeneity




Completeness




V Measure




Adjusted Rand




Adj. Mutual Info



Number of Clusters



Best Epsilon = 0.9  
Best N\_Samples = 50.0



Estimated Number of Clusters: 7.0

Best DBSCAN Scores for Cosine Distance Metric

Homogeneity score: 0.7896799513244787

Completeness score: 0.667782538810624

V\_measure score: 0.7236336939202158

Adjusted Rand Index score: 0.7108732847743878

Adjusted Mutual Information Score: 0.7221232251101021

```
[152]: # Find best parameter combo in DBSCAN with EUCLIDEAN metric
# Generate 3D surface plots for each metric against epsilon and min samples
→values

umap_reducer = umap.UMAP(n_components = 5, init = 'random', metric =
→'euclidean', random_state=0)
bbc_umap_EUC = umap_reducer.fit_transform(bbc_X_tfidf)

eps_vals = [.01,.05,.1,.2,.3,.4,.5,.6,.7,.8,.9,1]
```

```

samples = [1,3,5,10,20,30,35,40,50]
metric_scores_DBSCAN_EUC = []

for eps in eps_vals:
    for sam in samples:
        DBSCAN_clust = DBSCAN(min_samples=sam, eps=eps)

        DBSCAN_clust.fit(bbc_umap)

        labels = DBSCAN_clust.labels_
        n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)

        metric_scores_DBSCAN_EUC.append([metrics.homogeneity_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.completeness_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.v_measure_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.adjusted_rand_score(bbc.target,
→DBSCAN_clust.labels_),
                                         metrics.adjusted_mutual_info_score(bbc.
→target, DBSCAN_clust.labels_),
                                         n_clusters_,eps,sam])

# Visualize parameter effects on metrics
from mpl_toolkits.mplot3d import Axes3D
from scipy.interpolate import griddata
metricdata = np.array(metric_scores_DBSCAN_EUC)
metric_scores = ['Homogeneity', 'Completeness', 'V Measure', 'Adjusted Rand',
→'Adj. Mutual Info', 'Number of Clusters']

for i in range(0,6):
    plotdata = metricdata[:, [i,6,7]]
    z, x, y = zip(*plotdata)
    z = list(map(float, z))
    grid_x, grid_y = np.mgrid[min(x):max(x):100j, min(y):max(y):100j]
    grid_z = griddata((x, y), z, (grid_x, grid_y), method='cubic')

    fig = plt.figure()
    ax = fig.gca(projection='3d')
    ax.plot_surface(grid_x, grid_y, grid_z, cmap=plt.cm.Spectral)
    ax.set_xlabel('Epsilon')
    ax.set_ylabel('Min. Samples')
    ax.set_zlabel('Score')
    ax.set_title(str(metric_scores[i]))

```

```


plt.show()

# Find best parameter combination for epsilon and minimum samples


avs = []
for i in range(len(metricdata)):
    a = ␣
    →((metricdata[i][0]+metricdata[i][1]+metricdata[i][2]+metricdata[i][3]+metricdata[i][4])/
    →5)
    avs.append(a)

b = avs.index(max(avs))
print('Best Epsilon = ', metricdata[b][6])
print('Best N_Samples = ', metricdata[b][7])
print()
cm = contingency_matrix(bbc.target, DBSCAN_clust.labels_)
rows, cols=linear_sum_assignment(cm, maximize=True)
plot_mat(cm[rows[:, np.newaxis], cols], xticklabels=cols, yticklabels=rows,␣
→size=(5,5))
print("Estimated Number of Clusters: ", metricdata[b][5])
print()
print("Best DBSCAN Scores for Euclidean Distance Metric")
print("Homogeneity score: ", metricdata[b][0])
print("Completeness score: ", metricdata[b][1])
print("V_measure score: ", metricdata[b][2])
print("Adjusted Rand Index score: ", metricdata[b][3])
print("Adjusted Mutual Information Score: ", metricdata[b][4])


```




Completeness




V Measure



Adjusted Rand




Adj. Mutual Info






### Number of Clusters



Best Epsilon = 0.9  
Best N\_Samples = 50.0



Estimated Number of Clusters: 7.0

Best DBSCAN Scores for Euclidean Distance Metric

Homogeneity score: 0.7896799513244787

Completeness score: 0.667782538810624

V\_measure score: 0.7236336939202158

Adjusted Rand Index score: 0.7108732847743878

Adjusted Mutual Information Score: 0.7221232251101021

[ ]: